

Lab Assignment 3

Author: Vendramin Nicolò

In order to solve the problem presented in task one I had to develop a filter function and a map function.

The filter function is required to only keep in the streams those ride events that are starting (or ending, depending whether it is a start or end event) in one of the terminals of the JFK airport of New York.

The mapper instead is used to put the stream from a sequence of TaxiRides to a sequence of tuples showing <Terminal_String, Cardinality, Hour of the day>. Below I attach the code of the two classes:

```
/*
Filters only the start events or end events in a JFKTerminal.
Keeps only those taxi rides that are start events or end events having, respectively as
a starting or ending location, one of the terminals of the JFK Airport.
*/
public static class JFKFilter implements FilterFunction<TaxiRide> {

    @Override
    public boolean filter(TaxiRide taxiRide) throws Exception {

        JFKTerminal terminal;

        // If the record is a start event
        if(taxiRide.isStart())
            // consider as location the starting location
            terminal = JFKTerminal.gridToTerminal(GeoUtils.mapToGridCell(taxiRide.startLon, taxiRide.startLat));
        // If the record is an end event
        else
            // consider as location the ending location
            terminal = JFKTerminal.gridToTerminal(GeoUtils.mapToGridCell(taxiRide.endLon, taxiRide.endLat));

        // the condition to filter is that the location is a terminal of JFK Airport
        boolean condition = terminal != JFKTerminal.NOT_A_TERMINAL;

        if(condition)
            return true;
        else return false;
    }
}
```

```
/*
This mapper maps each event to its Terminal, 1, hour tuple.
*/
public static final class TerminalPresenceTimeMapper implements MapFunction<TaxiRide, Tuple3<JFKTerminal, Integer, Integer>> {

    @Override
    public Tuple3<JFKTerminal, Integer, Integer> map(TaxiRide taxiRide) throws Exception {

        int grid = 0;
        long millis = 0;

        // If the record is a start event
        if(taxiRide.isStart()) {
            // the cell is extracted from the starting location
            grid = GeoUtils.mapToGridCell(taxiRide.startLon, taxiRide.startLat);
            // the time is extracted from the starting time
            millis = taxiRide.startTime.getMillis();
        }
        // If the record is an end event
        else{
            // the cell is extracted from the end location
            grid = GeoUtils.mapToGridCell(taxiRide.endLon, taxiRide.endLat);
            // the time is extracted from the end time
            millis = taxiRide.endTime.getMillis();
        }

        // We set up a calendar to be able to extract the hour of the day basing on the unix timestamp
        Calendar calendar = Calendar.getInstance();
        calendar.setTimeZone(TimeZone.getTimeZone( "America/New_York" ));
        calendar.setTimeInMillis(millis);

        // We return a tuple including the terminal of the record, the number of events (1), and the hour of the day
        return new Tuple3<>(JFKTerminal.gridToTerminal(grid), value: 1, calendar.get(Calendar.HOUR_OF_DAY));
    }
}
```

For what the second task is concerned I was having two possible interpretations. For one of them I needed to implement a second filter that only keeps those ride records corresponding to start event. Below I attach the code:

```
/*
Filters only those rides the start ride events
*/
public static class StartRideFilter implements FilterFunction<TaxiRide> {

    @Override
    public boolean filter(TaxiRide taxiRide) throws Exception {

        return taxiRide.isStart;
    }
}
```

In the main() I proceed as follow:

At first the execution environment is set up, and the file containing the data is added as a source to the execution environment.

The watermark is set to 60 seconds, considering possible disconnections or delay in the communication of the events from the car to the system storing the data.

```
StreamExecutionEnvironment env = StreamExecutionEnvironment.getExecutionEnvironment();
env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

// get the taxi ride data stream form the file
DataStream<TaxiRide> rides = env.addSource(
    new TaxiRideSource(
        dataFilePath: "/Users/nicolovendramin/flinkLab/flink-java-project/src/main/" +
            "java/org/apache/flink/quickstart/data/nycTaxiRides.gz",
        maxEventDelaySecs: 60, // Watermark
        servingSpeedFactor: 2000));
```

After initialising the data stream we can start applying the transformations for each task.

Task1: For each hour progressively, for each terminal count the number of events (both start and end events).

```
// Generating the result of task1
DataStream<Tuple3<JFKTerminal, Integer, Integer>> terminal_rides = rides
    .filter(new JFKFilter()) // filtering rides starting or arriving in JFK terminals
    .map(new TerminalPresenceTimeMapper()) // mapping each ride to the leave-arrive events
    .keyBy( ...fields: 2) // grouping the result by hour of the day
    .keyBy( ...fields: 0) // grouping inside each single grouping by terminal
    .timeWindow(Time.hours(1)) // defining a one hour time window
    .sum( positionToSum: 1); // summing over the grouping in the desired time window
```

Simply we filter only the rides starting or ending at JFK airport, we map them to the <Terminal_String, Cardinality, Hour of the day> format using the map class reported above, we group them by hour of the day and by terminal (in this order), we define a timeWindow of 1 hour and we perform a sum over the elements of the grouping in the required time window.

Output(Task1):

```
7> (TERMINAL_5,3,0)
4> (TERMINAL_6,9,1)
2> (TERMINAL_4,70,1)
7> (TERMINAL_3,28,1)
3> (TERMINAL_1,1,1)
6> (TERMINAL_2,8,1)
2> (TERMINAL_4,47,2)
7> (TERMINAL_3,48,2)
7> (TERMINAL_5,1,2)
3> (TERMINAL_1,2,2)
4> (TERMINAL_6,12,2)
6> (TERMINAL_2,3,2)
2> (TERMINAL_4,42,3)
4> (TERMINAL_6,9,3)
6> (TERMINAL_2,1,3)
```

Task2(alpha): For each hour of the day we pick the busiest terminal with the number of events (both start and end events).

```
// Generating the result of task2 from previous exercise
DataStream<Tuple3<JFKTerminal, Integer, Integer>> terminal_rides_max = terminal_rides
    .keyBy( ...fields: 2) // grouping the previous result by hour of the day
    .timeWindowAll(Time.hours(1)) // selecting a time window of one hour across all nodes
    .max( positionToMax: 1); // picking the max with respect to the counting
```

We just start from the previously obtained stream, we apply a further grouping by hour of the day, we set a time window across all the nodes and we extract the maximum over the different nodes.

Output(Task2(alpha)):

```
7> (TERMINAL_6,70,1)
8> (TERMINAL_6,48,2)
1> (TERMINAL_6,42,3)
2> (TERMINAL_6,45,4)
3> (TERMINAL_1,31,5)
4> (TERMINAL_1,26,6)
5> (TERMINAL_6,50,7)
6> (TERMINAL_6,57,8)
7> (TERMINAL_1,73,9)
```

Task2(beta): For each hour of the day we pick the terminal from which more people has to leave (only start events).

```
// Generating the result of task2 considering only trips leaving the terminal
DataStream<Tuple3<JFKTerminal, Integer, Integer>> terminal_rides_max_leaving = rides
    .filter(new StartRideFilter()) // filtering on start rides
    .filter(new JFKFilter()) // filtering rides leaving from JFK terminals
    .map(new TerminalPresenceTimeMapper()) // mapping each ride to the leave events
    .keyBy( ...fields: 2) // grouping the result by hour of the day
    .keyBy( ...fields: 0) // grouping inside each single grouping by terminal
    .timeWindow(Time.hours(1)) // defining a one hour time window
    .sum( positionToSum: 1) // summing over the grouping in the desired time window
    .keyBy( ...fields: 2) // grouping the previous result by hour of the day
    .timeWindowAll(Time.hours(1)) // selecting a time window of one hour across all nodes
    .max( positionToMax: 1); // picking the max with respect to the counting
```

Here we have to restart from the clean. (imported from file) data stream. We apply a first filter to keep only the start event records. After that we perform all the operations that have been described for task 1 and task2(alpha) in the same order.

Output(Task2(beta)):

```
5> (TERMINAL_2,10,23)
6> (TERMINAL_6,43,0)
7> (TERMINAL_3,63,1)
8> (TERMINAL_4,46,2)
1> (TERMINAL_2,41,3)
2> (TERMINAL_3,41,4)
3> (TERMINAL_2,26,5)
4> (TERMINAL_6,20,6)
5> (TERMINAL_2,44,7)
6> (TERMINAL_3,54,8)
7> (TERMINAL_1,63,9)
8> (TERMINAL_6,45,10)
```

The full code is attached at the end of this document in case of any doubt about how the different pieces are plugged.

```

HomeworkRight.java x
1 // Just importing all the necessary namespaces
2 package org.apache.flink.quickstart;
3 import com.dataartisans.flinktraining.exercises.datastream_java.datatypes.TaxiRide;
4 import com.dataartisans.flinktraining.exercises.datastream_java.sources.TaxiRideSource;
5 import com.dataartisans.flinktraining.exercises.datastream_java.utils.GeoUtils;
6 import org.apache.flink.api.common.functions.FilterFunction;
7 import org.apache.flink.api.common.functions.MapFunction;
8 import org.apache.flink.api.java.tuple.Tuple3;
9 import org.apache.flink.streaming.api.TimeCharacteristic;
10 import org.apache.flink.streaming.api.datastream.DataStream;
11 import org.apache.flink.streaming.api.environment.StreamExecutionEnvironment;
12 import org.apache.flink.streaming.api.windowing.time.Time;
13 import java.util.Scanner;
14 import java.util.Calendar;
15 import java.util.TimeZone;
16
17
18 public class HomeworkRight{
19
20     // Already provided util to handle terminals and their location
21     public enum JFKTerminal {...}
22
23
24     /* Main body of the execution containing the data importation, the handling
25     of user choice of the task to execute, the stream processing logic and the
26     production of the output.
27     */
28     public static void main(String[] args) throws Exception {
29
30         StreamExecutionEnvironment env = StreamExecutionEnvironment.getExecutionEnvironment();
31         env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);
32
33         // get the taxi ride data stream from the file
34         DataStream<TaxiRide> rides = env.addSource(
35             new TaxiRideSource(
36                 dataFilePath: "/Users/nicolovendramin/flinkLab/flink-java-project/src/main/" +
37                     "java/org/apache/flink/quickstart/data/nycTaxiRides.gz",
38                 maxEventDelaySecs: 60, // Watermark
39                 servingSpeedFactor: 2000));
40
41         // Reading from System.in to know which one of the tasks we want to print.
42         Scanner reader = new Scanner(System.in);
43
44         // We keep reading until we get a valid choice
45         int n = -1;
46         while(n<0 || n>3) {
47             System.out.println(
48                 "Enter the number of the task you want to print " +
49                 "[1 -> terminal visit per hour," +
50                 " 2-> busiest terminal per hour," +
51                 " 3 -> busiest terminal in exit per hour," +
52                 " 0 -> all of them]: ");
53
54             n = reader.nextInt(); // Scans the next token of the input as an int.
55         }
56         reader.close();
57
58         // Generating the result of task1
59         DataStream<Tuple3<JFKTerminal, Integer, Integer>> terminal_rides = rides
60             .filter(new JFKFilter()) // filtering rides starting or arriving in JFK terminals
61             .map(new TerminalPresenceTimeMapper()) // mapping each ride to the leave-arrive events
62             .keyBy( ...fields: 2) // grouping the result by hour of the day
63             .keyBy( ...fields: 0) // grouping inside each single grouping by terminal
64             .timeWindow(Time.hours(1)) // defining a one hour time window
65             .sum( positionToSum: 1); // summing over the grouping in the desired time window
66
67         // Generating the result of task2 from previous exercise
68         DataStream<Tuple3<JFKTerminal, Integer, Integer>> terminal_rides_max = terminal_rides
69             .keyBy( ...fields: 2) // grouping the previous result by hour of the day
70             .timeWindowAll(Time.hours(1)) // selecting a time window of one hour across all nodes
71             .max( positionToMax: 1); // picking the max with respect to the counting
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92

```

```

93 // Generating the result of task2 considering only trips leaving the terminal
94 DataStream<Tuple3<JFKTerminal, Integer, Integer>> terminal_rides_max_leaving = rides
95     .filter(new StartRideFilter()) // filtering on start rides
96     .filter(new JFKFilter()) // filtering rides leaving from JFK terminals
97     .map(new TerminalPresenceTimeMapper()) // mapping each ride to the leave events
98     .keyBy(...fields: 2) // grouping the result by hour of the day
99     .keyBy(...fields: 0) // grouping inside each single grouping by terminal
100     .timeWindow(Time.hours(1)) // defining a one hour time window
101     .sum(positionToSum: 1) // summing over the grouping in the desired time window
102     .keyBy(...fields: 2) // grouping the previous result by hour of the day
103     .timeWindowAll(Time.hours(1)) // selecting a time window of one hour across all nodes
104     .max(positionToMax: 1); // picking the max with respect to the counting
105
106 // printing only the required results
107 if(n == 1 || n == 0){
108     terminal_rides.print();
109 }
110 if(n == 2 || n == 0) {
111     terminal_rides_max.print();
112 }
113 if(n == 3 || n == 0) {
114     terminal_rides_max_leaving.print();
115 }
116
117 env.execute();
118 }
119
120
121 /*
122 Filters only the start events or end events in a JFKTerminal.
123 Keeps only those taxi rides that are start events or end events having, respectively as
124 a starting or ending location, one of the terminals of the JFK Airport.
125 */
126 public static class JFKFilter implements FilterFunction<TaxiRide> {
127
128     @Override
129     public boolean filter(TaxiRide taxiRide) throws Exception {
130
131         JFKTerminal terminal;
132
133         // If the record is a start event
134         if(taxiRide.isStart)
135             // consider as location the starting location
136             terminal = JFKTerminal.gridToTerminal(GeoUtils.mapToGridCell(taxiRide.startLon, taxiRide.startLat));
137         // If the record is an end event
138         else
139             // consider as location the ending location
140             terminal = JFKTerminal.gridToTerminal(GeoUtils.mapToGridCell(taxiRide.endLon, taxiRide.endLat));
141
142         // the condition to filter is that the location is a terminal of JFK Airport
143         boolean condition = terminal != JFKTerminal.NOT_A_TERMINAL;
144
145         if(condition)
146             return true;
147         else return false;
148     }
149 }
150
151
152 /*
153 Filters only those rides the start ride events
154 */
155 public static class StartRideFilter implements FilterFunction<TaxiRide> {
156
157     @Override
158     public boolean filter(TaxiRide taxiRide) throws Exception {
159
160         return taxiRide.isStart;
161     }
162 }
163

```

```

93 // Generating the result of task2 considering only trips leaving the terminal
94 DataStream<Tuple3<JFKTerminal, Integer, Integer>> terminal_rides_max_leaving = rides
95     .filter(new StartRideFilter()) // filtering on start rides
96     .filter(new JFKFilter()) // filtering rides leaving from JFK terminals
97     .map(new TerminalPresenceTimeMapper()) // mapping each ride to the leave events
98     .keyBy( ...fields: 2) // grouping the result by hour of the day
99     .keyBy( ...fields: 0) // grouping inside each single grouping by terminal
100     .timeWindow(Time.hours(1)) // defining a one hour time window
101     .sum( positionToSum: 1) // summing over the grouping in the desired time window
102     .keyBy( ...fields: 2) // grouping the previous result by hour of the day
103     .timeWindowAll(Time.hours(1)) // selecting a time window of one hour across all nodes
104     .max( positionToMax: 1); // picking the max with respect to the counting
105
106 // printing only the required results
107 if(n == 1 || n == 0){
108     terminal_rides.print();
109 }
110 if(n == 2 || n == 0) {
111     terminal_rides_max.print();
112 }
113 if(n == 3 || n == 0) {
114     terminal_rides_max_leaving.print();
115 }
116
117 env.execute();
118
119 }
120
121 /*
122 Filters only the start events or end events in a JFKTerminal.
123 Keeps only those taxi rides that are start events or end events having, respectively as
124 a starting or ending location, one of the terminals of the JFK Airport.
125 */
126 public static class JFKFilter implements FilterFunction<TaxiRide> {
127
128     @Override
129     public boolean filter(TaxiRide taxiRide) throws Exception {
130
131         JFKTerminal terminal;
132
133         // If the record is a start event
134         if(taxiRide.isStart)
135             // consider as location the starting location
136             terminal = JFKTerminal.gridToTerminal(GeoUtils.mapToGridCell(taxiRide.startLon, taxiRide.startLat));
137         // If the record is an end event
138         else
139             // consider as location the ending location
140             terminal = JFKTerminal.gridToTerminal(GeoUtils.mapToGridCell(taxiRide.endLon, taxiRide.endLat));
141
142         // the condition to filter is that the location is a terminal of JFK Airport
143         boolean condition = terminal != JFKTerminal.NOT_A_TERMINAL;
144
145         if(condition)
146             return true;
147         else return false;
148     }
149 }
150
151 /*
152 Filters only those rides the start ride events
153 */
154 public static class StartRideFilter implements FilterFunction<TaxiRide> {
155
156     @Override
157     public boolean filter(TaxiRide taxiRide) throws Exception {
158
159         return taxiRide.isStart;
160     }
161 }
162
163

```

```

164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204

```

```

/*
This mapper maps each event to its Terminal, 1, hour tuple.
*/
public static final class TerminalPresenceTimeMapper implements MapFunction<TaxiRide, Tuple3<JFKTerminal, Integer, Integer>> {

    @Override
    public Tuple3<JFKTerminal, Integer, Integer> map(TaxiRide taxiRide) throws Exception {

        int grid = 0;
        long millis = 0;

        // If the record is a start event
        if(taxiRide.isStart()) {
            // the cell is extracted from the starting location
            grid = GeoUtils.mapToGridCell(taxiRide.startLon, taxiRide.startLat);
            // the time is extracted from the starting time
            millis = taxiRide.startTime.getMillis();
        }
        // If the record is an end event
        else{
            // the cell is extracted from the end location
            grid = GeoUtils.mapToGridCell(taxiRide.endLon, taxiRide.endLat);
            // the time is extracted from the end time
            millis = taxiRide.endTime.getMillis();
        }

        // We set up a calendar to be able to extract the hour of the day basing on the unix timestamp
        Calendar calendar = Calendar.getInstance();
        calendar.setTimeZone(TimeZone.getTimeZone( "America/New_York" ));
        calendar.setTimeInMillis(millis);

        // We return a tuple including the terminal of the record, the number of events (1), and the hour of the day
        return new Tuple3<>(JFKTerminal.gridToTerminal(grid), value: 1, calendar.get(Calendar.HOUR_OF_DAY));
    }
}

```


